

```

!git clone https://bitbucket.org/jadslim/german-traffic-signs

Cloning into 'german-traffic-signs'...
Unpacking objects: 100% (6/6), 117.80 MiB | 8.64 MiB/s, done.
Updating files: 100% (4/4), done.

!ls german-traffic-sign

ls: cannot access 'german-traffic-sign': No such file or directory

import numpy as np
import matplotlib.pyplot as plt
import keras
from keras.models import Sequential
from keras.layers import Dense
from keras.optimizers import Adam
from keras.utils import to_categorical
from keras.layers import Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
import random
import pickle
import pandas as pd
import cv2
from keras.callbacks import LearningRateScheduler, ModelCheckpoint

%matplotlib inline

np.random.seed(0)

# TODO: Implement load the data here.
with open('german-traffic-signs/train.p', 'rb') as f:
    train_data = pickle.load(f)
with open('german-traffic-signs/valid.p', 'rb') as f:
    val_data = pickle.load(f)
# TODO: Load test data
with open('german-traffic-signs/test.p', 'rb') as f:
    test_data = pickle.load(f)

# Split out features and labels
X_train, y_train = train_data['features'], train_data['labels']
X_val, y_val = val_data['features'], val_data['labels']
X_test, y_test = test_data['features'], test_data['labels']

# 4 dimensional
print(X_train.shape)
print(X_test.shape)
print(X_val.shape)

```

```

(34799, 32, 32, 3)
(12630, 32, 32, 3)
(4410, 32, 32, 3)

assert(X_train.shape[0] == y_train.shape[0]), "The number of images is
not equal to the number of labels."
assert(X_train.shape[1:] == (32,32,3)), "The dimensions of the images
are not 32 x 32 x 3."
assert(X_val.shape[0] == y_val.shape[0]), "The number of images is not
equal to the number of labels."
assert(X_val.shape[1:] == (32,32,3)), "The dimensions of the images
are not 32 x 32 x 3."
assert(X_test.shape[0] == y_test.shape[0]), "The number of images is
not equal to the number of labels."
assert(X_test.shape[1:] == (32,32,3)), "The dimensions of the images
are not 32 x 32 x 3."

data = pd.read_csv('german-traffic-signs/signnames.csv')

num_of_samples=[]

cols = 5
num_classes = 43

fig, axs = plt.subplots(nrows=num_classes, ncols=cols, figsize=(5,50))
fig.tight_layout()

for i in range(cols):
    for j, row in data.iterrows():
        x_selected = X_train[y_train == j]
        axs[j][i].imshow(x_selected[random.randint(0,(len(x_selected)
- 1)), :, :], cmap=plt.get_cmap('gray'))
        axs[j][i].axis("off")
        if i == 2:
            axs[j][i].set_title(str(j) + " - " + row["SignName"])
            num_of_samples.append(len(x_selected))

```

0 - Speed limit (20km/h)



1 - Speed limit (30km/h)



2 - Speed limit (50km/h)



3 - Speed limit (60km/h)



4 - Speed limit (70km/h)



5 - Speed limit (80km/h)



6 - End of speed limit (80km/h)



7 - Speed limit (100km/h)



```


print(num_of_samples)
plt.figure(figsize=(12, 4))
plt.bar(range(0, num_classes), num_of_samples)
plt.title("Distribution of the train dataset")
plt.xlabel("Class number")
plt.ylabel("Number of images")
plt.show()

```

```

[180, 1980, 2010, 1260, 1770, 1650, 360, 1290, 1260, 1320, 1800, 1170,
1890, 1920, 690, 540, 360, 990, 1080, 180, 300, 270, 330, 450, 240,
1350, 540, 210, 480, 240, 390, 690, 210, 599, 360, 1080, 330, 180,
1860, 270, 300, 210, 210]

```



```

import cv2


plt.imshow(X_train[1000])
plt.axis("off")
print(X_train[1000].shape)
print(y_train[1000])

(32, 32, 3)
36


```




```
def grayscale(img):  
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)  
    return img  
img = grayscale(X_train[1000])  
plt.imshow(img)  
plt.axis("off")  
print(img.shape)  
(32, 32)
```



```
def equalize(img):  
    img = cv2.equalizeHist(img)  
    return img  
img = equalize(img)  
plt.imshow(img)  
plt.axis("off")  
print(img.shape)  
(32, 32)
```



```
def preprocess(img):  
    img = grayscale(img)  
    img = equalize(img)  
    img = img/255  
    return img  
  
X_train = np.array(list(map(preprocess, X_train)))  
X_test = np.array(list(map(preprocess, X_test)))  
X_val = np.array(list(map(preprocess, X_val)))  
  
plt.imshow(X_train[random.randint(0, len(X_train) - 1)])  
plt.axis('off')  
print(X_train.shape)  
  
(34799, 32, 32)
```



```
X_train = X_train.reshape(34799, 32, 32, 1)
X_test = X_test.reshape(12630, 32, 32, 1)
X_val = X_val.reshape(4410, 32, 32, 1)

from keras.preprocessing.image import ImageDataGenerator

datagen = ImageDataGenerator(width_shift_range=0.1,
                             height_shift_range=0.1,
                             zoom_range=0.2,
                             shear_range=0.1,
                             rotation_range=10.)

datagen.fit(X_train)

# for X_batch, y_batch in
batches = datagen.flow(X_train, y_train, batch_size = 15)
X_batch, y_batch = next(batches)


fig, axs = plt.subplots(1, 15, figsize=(20, 5))
fig.tight_layout()

for i in range(15):
    axs[i].imshow(X_batch[i].reshape(32, 32))
    axs[i].axis("off")

print(X_batch.shape)
```



(15, 32, 32, 1)



```
y_train = to_categorical(y_train, 43)
y_test = to_categorical(y_test, 43)
y_val = to_categorical(y_val, 43)

# create model

def modified_model():
    model = Sequential()
    model.add(Conv2D(60, (5, 5), input_shape=(32, 32, 1),
activation='relu'))
    model.add(Conv2D(60, (5, 5), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(30, (3, 3), activation='relu'))
    model.add(Conv2D(30, (3, 3), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Flatten())
    model.add(Dense(500, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(43, activation='softmax'))

    model.compile(Adam(lr = 0.001), loss='categorical_crossentropy',
metrics=['accuracy'])
    return model

model = modified_model()
print(model.summary())

history = model.fit_generator(datagen.flow(X_train, y_train,
batch_size=50),
                                steps_per_epoch=500,
                                epochs=20,
                                validation_data=(X_val, y_val), shuffle =
1)
```

WARNING:absl:`lr` is deprecated in Keras optimizer, please use  
`learning\_rate` or use the legacy optimizer,  
e.g.,tf.keras.optimizers.legacy.Adam.

Model: "sequential"


Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 60)	1560

conv2d_1 (Conv2D)	(None, 24, 24, 60)	90060
max_pooling2d (MaxPooling2D)	(None, 12, 12, 60)	0
conv2d_2 (Conv2D)	(None, 10, 10, 30)	16230
conv2d_3 (Conv2D)	(None, 8, 8, 30)	8130
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 30)	0
flatten (Flatten)	(None, 480)	0
dense (Dense)	(None, 500)	240500
dropout (Dropout)	(None, 500)	0
dense_1 (Dense)	(None, 43)	21543
=====		
Total params: 378023 (1.44 MB)		
Trainable params: 378023 (1.44 MB)		
Non-trainable params: 0 (0.00 Byte)		
=====		
<ipython-input-23-f90eee302bcd>:4: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.		
history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=50),		
None		
Epoch 1/20		
500/500 [=====] - 24s 25ms/step - loss: 2.0057 - accuracy: 0.4365 - val_loss: 0.4613 - val_accuracy: 0.8533		
Epoch 2/20		
500/500 [=====] - 12s 24ms/step - loss: 0.6879 - accuracy: 0.7870 - val_loss: 0.1464 - val_accuracy: 0.9587		
Epoch 3/20		
500/500 [=====] - 12s 25ms/step - loss: 0.4144 - accuracy: 0.8722 - val_loss: 0.1113 - val_accuracy: 0.9694		
Epoch 4/20		
500/500 [=====] - 11s 23ms/step - loss: 0.3062 - accuracy: 0.9069 - val_loss: 0.0807 - val_accuracy: 0.9780		
Epoch 5/20		
500/500 [=====] - 11s 22ms/step - loss: 0.2400 - accuracy: 0.9270 - val_loss: 0.0967 - val_accuracy: 0.9721		
Epoch 6/20		
500/500 [=====] - 12s 24ms/step - loss:		

0.2047 - accuracy: 0.9358 - val\_loss: 0.0709 - val\_accuracy: 0.9798  
Epoch 7/20  
500/500 [=====] - 12s 24ms/step - loss:  
0.1886 - accuracy: 0.9403 - val\_loss: 0.0453 - val\_accuracy: 0.9882  
Epoch 8/20  
500/500 [=====] - 12s 24ms/step - loss:  
0.1643 - accuracy: 0.9498 - val\_loss: 0.0457 - val\_accuracy: 0.9864  
Epoch 9/20  
500/500 [=====] - 11s 23ms/step - loss:  
0.1516 - accuracy: 0.9520 - val\_loss: 0.0462 - val\_accuracy: 0.9891  
Epoch 10/20  
500/500 [=====] - 12s 23ms/step - loss:  
0.1394 - accuracy: 0.9583 - val\_loss: 0.0480 - val\_accuracy: 0.9891  
Epoch 11/20  
500/500 [=====] - 12s 23ms/step - loss:  
0.1247 - accuracy: 0.9601 - val\_loss: 0.0464 - val\_accuracy: 0.9864  
Epoch 12/20  
500/500 [=====] - 12s 23ms/step - loss:  
0.1197 - accuracy: 0.9630 - val\_loss: 0.0365 - val\_accuracy: 0.9893  
Epoch 13/20  
500/500 [=====] - 12s 23ms/step - loss:  
0.1102 - accuracy: 0.9663 - val\_loss: 0.0423 - val\_accuracy: 0.9866  
Epoch 14/20  
500/500 [=====] - 11s 23ms/step - loss:  
0.1056 - accuracy: 0.9669 - val\_loss: 0.0499 - val\_accuracy: 0.9857  
Epoch 15/20  
500/500 [=====] - 11s 23ms/step - loss:  
0.1066 - accuracy: 0.9684 - val\_loss: 0.0389 - val\_accuracy: 0.9884  
Epoch 16/20  
500/500 [=====] - 11s 22ms/step - loss:  
0.0983 - accuracy: 0.9708 - val\_loss: 0.0283 - val\_accuracy: 0.9932  
Epoch 17/20  
500/500 [=====] - 11s 22ms/step - loss:  
0.0971 - accuracy: 0.9692 - val\_loss: 0.0278 - val\_accuracy: 0.9943  
Epoch 18/20  
500/500 [=====] - 11s 23ms/step - loss:  
0.0817 - accuracy: 0.9756 - val\_loss: 0.0229 - val\_accuracy: 0.9941  
Epoch 19/20  
500/500 [=====] - 12s 23ms/step - loss:  
0.0859 - accuracy: 0.9732 - val\_loss: 0.0352 - val\_accuracy: 0.9900  
Epoch 20/20  
500/500 [=====] - 12s 24ms/step - loss:  
0.0889 - accuracy: 0.9732 - val\_loss: 0.0300 - val\_accuracy: 0.9941


```
plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.title('Loss')  
plt.xlabel('epoch')
```

```
Text(0.5, 0, 'epoch')
```



```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'test'])
plt.title('Accuracy')
plt.xlabel('epoch')

# TODO: Evaluate model on test data
score = model.evaluate(X_test, y_test, verbose=0)
```



```
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
Test score: 0.13481435179710388
Test accuracy: 0.9684877395629883
```

```
#fetch image
```

```
import requests
from PIL import Image
url = 'https://c8.alamy.com/comp/J2MRAJ/german-road-sign-bicycles-
crossing-J2MRAJ.jpg'
r = requests.get(url, stream=True)
img = Image.open(r.raw)
plt.imshow(img, cmap=plt.get_cmap('gray'))
```

```
#Preprocess image
```

```
img = np.asarray(img)
img = cv2.resize(img, (32, 32))
img = preprocess(img)
plt.imshow(img, cmap = plt.get_cmap('gray'))
print(img.shape)
```

```
#Reshape reshape
```

```
img = img.reshape(1, 32, 32, 1)
```

```
#Test image
```

```
print("predicted sign: "+ str(np.argmax(model.predict(img), axis=-1)))
```

```
(32, 32)
```

```
1/1 [=====] - 0s 18ms/step
```

```
predicted sign: [29]
```

